

Linking

- Linking and Libraries
- Dynamic Linking Tradeoffs
- Dynamic (runtime) Linking Interface
- Library Naming Conventions
- Library Versions
- Using Libraries
- Dynamic Loaded Expression Evaluator
- `dlexpr.h` Header File Selections
- `dlexpr.c` Auxiliary Routines
- `dlexpr.c` Auxiliary Routines Continued
- `dlexpr.c` `main()`
- `dlexpr.c` `main()` Continued
- Arithmetic Grammar
- Arithmetic Evaluation Routines
- Arithmetic Evaluation Routines Continued
- Arithmetic Evaluation Routines Continued
- References

Linking and Libraries

- A relocatable object module (typically with extension `.o`) is produced from a single `.c` file (and multiple header files) using the `-c` compiler option.
- Object file modules can be in different formats: a `.out` is a historical format; *Executable Linking Format* ELF is a more modern format (used on both Solaris and Linux).
- A library is a collection of object modules along with some kind of symbol table documenting the symbols imported and exported by each module.
- When a library is linked into an executable, only those object modules necessary to satisfy outstanding symbols are added to the executable; hence the order of searching of libraries is very important.

For example, to use a non-standard `malloc` library, it should be searched before the standard `libc` library.

- Static linking searches libraries at compile-time and builds a self-contained executable. Built by a compile-time *link editor* (`ld(1)` on Solaris).
- Dynamic-linking searches libraries at load-time and even runtime. Uses *runtime linker* (`ld.so.1(1)` on Solaris).
- A dynamically linked executable can link to a routine on a *need-to* basis at runtime. This is also referred to as *lazy linking*.

Dynamic Linking Tradeoffs

- Shared libraries are *shared*: different programs share the same code, thus using less memory.
- Executables depending on shared libraries do not contain the shared library code and thus are smaller.
- Shared library bugs can be fixed by providing a new version of the library, without needing to recompile/relink the programs dependent on it.
- Shared libraries increase complexity. A executable is not standalone but requires a specific environment.
- Allows interposition of *wrapper functions* for library functions; useful for tracing and profiling.

Dynamic (runtime) Linking Interface

```
void *dlopen (const char *fileName, int flag);  
void *dlsym(void *dlHandle, char *symName);  
const char *dlerror(void);  
int dlclose (void *dlHandle);
```

- `dlopen()` adds dynamic library in `fileName` to executable, and returns a opaque handle. `flag` may be `RTLD_LAZY` or `RTLD_NOW`. If `RTLD_GLOBAL` is or'ed with `flag`, then symbols in `fileName` are made available to subsequently loaded modules.
- `dlsym()` returns the address of `symName` in dynamic library referenced by `dlHandle`. `NULL` if not found.
- `dlerror()` returns a human-readable error message for last error.
- `dlclose()` closes a dynamic library and unloads it if its reference count has gone to 0.

Library Naming Conventions

- Library names start with prefix `lib` as in `libc`. The prefix is omitted when the library is specified on the compiler command-line: `-lc` specifies linking with `libc`.
- Extension `.a` is used for static libraries (historically produced using the `ar` program).
- Extension `.so` is used for dynamic shared libraries.
- Above conventions need not be followed by runtime libraries loaded by `dlopen()`.

Library Versions

The following conventions are used to keep track of library versions:

- Each shared library has a *soname* of the form *name* . *so* . *V* where *V* is the version number.
- When the library API is no longer backward compatible, the version number should change.
- Each library has a actual name of the form *name* . *so* . *V* . *M* . *R*, where *M* is a minor version number (no interface change) and *R* is a release number (may be a bug fix).
- When a library is created, the *soname* is specified on the linker command-line (along with the actual name).
- Once again, when linking against a library, it is the *soname* which is specified.
- The *soname*'s are symbolic links to the latest actual library name, thus enabling a seamless upgrade when the API is backward compatible.

Using Libraries

- The environmental variable `LD_LIBRARY_PATH` is a : separated list of directories searched by the linker for dynamic libraries.
- `ldd` lists the libraries a executable depends on:

```
$ ldd /bin/ls
libc.so.1 => /usr/lib/libc.so.1
libdl.so.1 => /usr/lib/libdl.so.1
/usr/platform/SUNW,Ultra-Enterprise/lib/libc_psr.so.1
$
```

Can be a security problem with untrusted executables.

Dynamic Loaded Expression Evaluator

- Program accepts a single argument, which is the name of dynamically loaded module which evaluates expressions of a certain kind.

```
$ LD_LIBRARY_PATH=. ./dlexpr arith
>> 1 + 3 * ( 2 + -4/3)
4
>> 1 + 3 * ( 2 + -4/3) *
ERR ^
>> $
$ LD_LIBRARY_PATH=. ./dlexpr bool
>> 0 | (1 & ~0)
1
>> 0 | (1 & ~0 | )
ERR ^
>> $
```

- Expression evaluators are written using recursive-descent techniques. `set jmp ()/long jmp ()` invaluable for error recovery and reporting.

dlexpr.h Header File Selections

```
typedef struct DLExprEnv {
    char ch;                /* lookahead */
    jmp_buf *jmpBufP;       /* error recovery */
    char line[LINE_MAX];    /* input line */
    unsigned lineIndex;     /* next char index */
    void(*nextCh)           /* function to get */
        (struct DLExprEnv *); /* next char */
} DLExprEnv;

/* A parse function takes a pointer to a DLExprEnv
 * and returns the value of the expression.
 */
typedef int (*ParseFn)(DLExprEnv *envP);

/* called if an error is encountered */
void error(DLExprEnv *envP);

/* Checks if lookahead matches specified char c.
 * If matches, then lookahead is advanced, else error.
 */
#define MATCH(c, envP) \
do { \
    if ((envP)->ch == (c)) { \
        (envP)->nextCh(envP); \
    } \
    else { \
        error(envP); \
    } \
} while (0)
```

dlexpr.c Auxiliary Routines

```
/* Prints a ^ char at error position.
 * aborts via longjmp().
 */
void
error(DLEExprEnv *envP)
{
    fprintf(stderr, "ERR%s^\n",
            envP->lineIndex - (envP->lineIndex > 0),
            "");
    longjmp(envP->jmpBufP, 1);
}

/* Update ch field of envP to next non-blank,
 * non-tab char from envP->line[].
 */
static void
nextChar(DLEExprEnv *envP)
{
    do {
        envP->ch = envP->line[envP->lineIndex++];
    } while (envP->ch == ' ' || envP->ch == '\t');
}
```

dlexpr.c Auxiliary Routines Continued

```
/* Read next line for stdin into envP->line[].
 * Returns 0 on EOF; non-zero otherwise.
 */
static int
getLine(DLEExprEnv *envP)
{
    while (1) {
        if (!fgets(envP->line, LINE_MAX, stdin)) {
            if (feof(stdin)) {
                return 0;
            }
            else {
                fprintf(stderr, "ERR I/O error\n");
            }
        }
        else if (envP->line[strlen(envP->line) - 1]
                 != '\n') {
            fprintf(stderr, "ERR Line too long\n");
        }
        else break;
    }
    envP->lineIndex = 0;
    envP->nextCh(envP);
    return 1;
}
```

dlexpr.c main()

```
#define PARSE_FN_NAME_SUFFIX          "Parse"

int
main(int argc, char *argv[])
{
    void *dlHandle;
    char *parseFnName;
    ParseFn parseFn;
    if (argc != 2) {
        fprintf(stderr, "usage: %s EXPR_MODULE\n",
            argv[0]);
        exit(1);
    }
    if (!(dlHandle = dlopen(argv[1], RTLD_NOW)) {
        fprintf(stderr, "%s dlopen(): %s\n",
            argv[0], dlerror());
        exit(1);
    }
    parseFnName =
        malloc(strlen(argv[1]) +
            strlen(PARSE_FN_NAME_SUFFIX) + 1);
    if (!parseFnName) {
        perror(argv[0]); exit(1);
    }
}
```

dlexpr.c main() Continued

```
strcpy(parseFnName, argv[1]);
strcpy(parseFnName + strlen(argv[1]),
        PARSE_FN_NAME_SUFFIX);
if (!(parseFn = dlsym(dlHandle, parseFnName))) {
    fprintf(stderr, "%s dlsym(): %s\n",
            argv[0], dlerror());
    exit(1);
}
{ /* The module is now loaded */
    jmp_buf jmpBuf;
    static DLEnv env;
    env.jmpBufP = &jmpBuf; env.nextCh = nextChar;
    if (setjmp(jmpBuf) >= 0) { //spec requires use within a test
        while (printf(">> ") && getLine(&env)) {
            printf("%d\n", parseFn(&env));
        }
    }
}

free(parseFnName);
dlclose(dlHandle);
return 0;
}
```

Arithmetic Grammar

```
arithParse
: exp '\n'
;

exp
: term { '+' term }
| term { '-' term }
| term
;

term
: factor { '*' factor }
| factor { '/' factor }
| factor
;

factor
: '(' exp ')'
| '-' factor
| DIGIT
;
```

Arithmetic Evaluation Routines

```
int
arithParse(DLEnv *envP)
{
    int v = expr(envP);
    MATCH('\n', envP);
    return v;
}

static int
expr(DLEnv *envP)
{
    int v = term(envP);
    while(envP->ch == '+' || envP->ch == '-') {
        char c = envP->ch;
        int v1;
        MATCH(c, envP);
        v1 = term(envP);
        v += (c == '+') ? v1 : -v1;
    }
    return v;
}
```

Arithmetic Evaluation Routines Continued

```
static int
term(DLEnv *envP)
{
    int v = factor(envP);
    while(envP->ch == '*' || envP->ch == '/') {
        char c = envP->ch;
        int v1;
        MATCH(c, envP);
        v1 = factor(envP);
        if (c == '*') {
            v *= v1;
        }
        else {
            v /= v1;
        }
    }
    return v;
}
```


Arithmetic Evaluation Routines Continued

```
static int
factor(DLEExprEnv *envP)
{
    int v;
    if (envP->ch == '(') {
        MATCH('(', envP);
        v = expr(envP);
        MATCH(')', envP);
    }
    else if (envP->ch == '-') {
        MATCH('-', envP);
        v = -factor(envP);
    }
    else if (isdigit(envP->ch)) {
        char c = envP->ch;
        MATCH(c, envP);
        v = c - '0';
    }
    else {
        error(envP);
    }
    return v;
}
```

References

Text, Ch. 41, 42.

David A. Wheeler, *Program Library HOWTO*, available online.

Oracle, *Linker and Libraries Guide*, available at <http://docs.oracle.com/cd/E19253-01/817-1984/>.

Michael K. Johnson and Erik W. Troan, *Linux Application Development*, Addison-Wesley, 1998, Ch. 7.

Greg Nakhimovsky, *Building library interposers for fun and profit*, Unix Insider, at <http://www.itworld.com/article/2798072/enterprise-software/building-library-interposers-for-fun-and-profit.html>.

Timothy W. Curry, *Profiling and Tracing Dynamic Library Usage via Interposition*, USENIX Conference Proceedings, Summer 1994, at <http://www.usenix.org/publications/library/proceedings/bos94/curry.html>.