

# Advanced I/O

- Overview
- Non-Blocking I/O
- Non-Blocking I/O on a FIFO
- Uses of Non-Blocking I/O on a FIFO
- Non-Blocking I/O Example
- Non-Blocking I/O Example Continued
- Non-Blocking I/O Example Continued
- Non-Blocking I/O Example Execution Log
- Synchronized I/O
- Integrity Measures
- Buffer Flushing System Calls
- `open()` Synchronization Flags
- Checking for I/O Ready Using `select()`
- `select()` Time Specifications
- Operating on File Descriptor Sets
- Select Example
- Select Example Continued
- Select Example Continued
- Controlling an Open File
- File Locking Using `fcntl()`
- Lock Inheritance
- Advisory versus Mandatory Locking

- Excluding Multiple Processes
- Excluding Multiple Processes Continued
- Excluding Multiple Processes Continued
- Excluding Multiple Processes Continued
- Excluding Multiple Processes Log
- Memory Mapped I/O
- Mapping Types
- Memory Mapped Files
- Memory Mapped I/O Continued
- Unmapping Memory
- Memory Mapped File Copy
- Memory Mapped File Copy Continued
- Shared Memory IPC
- `mprotect()` Call
- Auto Stack Growing
- Execution Log
- `main()` Program
- `main()` Program Continued
- Stack Initialization
- Stack Initialization Continued
- Stack Deallocation
- Stack Overflow Handler
- Postfix Interpreter
- Postfix Interpreter Continued

- References

# Overview

- Non-blocking I/O.
- Synchronous I/O.
- `select()` call.
- File locking.
- Memory-mapped I/O.

# Non-Blocking I/O

- Call to `open( )` can specify `O_NONBLOCK` flag.
- Use `fcntl( )` on a already open descriptor with `O_NONBLOCK` to make I/O non-blocking (must be used for pipes and sockets whose descriptors are not obtained using `open( )`).
- If a call would have blocked, then it returns with an error with `errno` set to `EAGAIN`.
- Affects I/O to terminals, pseudo-terminals, pipes, FIFOs, sockets. Can also affect I/O for regular files when mandatory locking enabled.

# Non-Blocking I/O on a FIFO

- `open( )` of FIFO with `O_NONBLOCK` for reading succeeds immediately irrespective of whether or not write end has been opened (it makes sense to read no data from a FIFO which is not open for writing).
- `open( )` of FIFO with `O_NONBLOCK` for writing succeeds immediately if read end has been opened; fails with error `ENXIO` if read end is not open (since writing to it would generate a `SIGPIPE` signal).
- Read on a non-blocking FIFO with no data fails with error `EAGAIN`.
- Write on a non-blocking FIFO which is full (cannot accomodate data without violating `PIPE_BUF` constraint) fails with error `EAGAIN`.

# Uses of Non-Blocking I/O on a FIFO

- Allows single process to open both ends of a FIFO by reading using `O_NONBLOCK`.
- Avoids deadlocks between processes opening two FIFOs: P1 opens FIFO A for reading and then FIFO B for writing; P2 opens FIFO B for reading and FIFO A for writing; deadlock without `O_NONBLOCK`.

# Non-Blocking I/O Example

```
int
main(int argc, char *argv[])
{
    int fifoFD;
    int fd[2];
    if (argc != 2) {
        fprintf(stderr, "usage: nonblock FIFO\n");
        exit(1);
    }
    if (fcntl(STDIN_FILENO, F_SETFL,
              O_NONBLOCK) < 0) {
        fprintf(stderr,
                "could not set STDIN non-block: %s\n",
                strerror(errno));
        exit(1);
    }
    if ((fifoFD = open(argv[1],
                       O_RDONLY|O_NONBLOCK)) < 0) {
        fprintf(stderr, "could not open %s: %s\n",
                argv[1], strerror(errno));
        exit(1);
    }
}
```



# Non-Blocking I/O Example Continued

```
fd[0] = STDIN_FILENO; fd[1] = fifoFD;
while (1) {
    int i;
    for (i = 0; i < 2; i++) {
        char *src =
            (i == 0) ? "STDIN" : "FIFO";
        enum { MAX_BUF = 128 };
        char buf[MAX_BUF];
        int n;
        if ((n = read(fd[i], buf, MAX_BUF))
            < 0) {
            if (errno != EAGAIN) {
                fprintf(stderr, "read error on "
                    "%s: %s\n",
                    src, strerror(errno));
                exit(1);
            }
        }
    }
}
```

# Non-Blocking I/O Example Continued

```
    else if (n == 0) {  
        return 0;  
    }  
    else {  
        printf("%s: %.*s", src, n, buf);  
    }  
    sleep(10);  
} /* for (i = 0; i < 2; i++) */  
} /* while (1) */  
}
```

# Non-Blocking I/O Example Execution Log

```
$ mkfifo fifo
$ ls -l fifo
prw-rw-r-- 1 umrigar umrigar          0 Mar 13 13:03 fifo
$ ./nonblock fifo
Hello from stdin
STDIN: Hello from stdin
FIFO: Hello from fifo
FIFO: Hello again from fifo
Hello again from stdin
STDIN: Hello again from stdin
$
```

Concurrently, in another shell:

```
$ echo "Hello from fifo" >fifo
$ echo "Hello again from fifo" >fifo
$
```

# Synchronized I/O

## *Synchronized*

Data is actually written to I/O device.

## *Synchronous*

Call returns only after data is read/written (usually, to buffer cache).

Normal Unix I/O is *synchronous* but *unsynchronized*. Asynchronous I/O uses a different group of system calls (`aio_write()`, etc).

# Integrity Measures

Need to distinguish between *data* (file contents) and *metadata* (information about the file like size and timestamps).

## **Data Integrity**

Data has been written out so that it can be retrieved. May require file size update but not necessarily file timestamps.

## **File Integrity**

Superset of data integrity: all data and meta-data has been written out.

# Buffer Flushing System Calls

```
void sync(void);  
int fsync(int fd);  
int fdatasync(int fd);
```

- `sync()` will flush the buffer cache at some future time (returns immediately). Used for shut down or unmounting removable device.
- `fsync()` does not return until **all buffers** associated with `fd` have been flushed to the I/O device. Ensures file integrity.
- `fdatasync()` does not return until **all data buffers** associated with `fd` have been flushed to the I/O device. Ensures data integrity.
- Since these are options, we need to check. If not supported, then `sync()` and `fsync()` may be NOP's!!

# `open ( )` Synchronization Flags

`O_SYNC`

Implicit `fsync ( )` after every write.

`O_DSYNC`

Implicit `fdatasync ( )` after every write.

`O_RSYNC`

Sync's reads. Must be used with `O_SYNC` or `O_DSYNC`. Causes inode access time to be flushed to disk. May disable read-ahead.

Using these flags can affect performance drastically.

# Checking for I/O Ready Using `select()`

```
int select(int maxfd1, fd_set *readfds,  
           fd_set *writefds,  
           fd_set *exceptfds,  
           struct timeval *tptr);
```

- `readfds` specifies a set of file descriptors on which to wait, until one or more is ready for reading.
- `writefds` specifies a set of file descriptors on which to wait, until one or more is ready for writing.
- `exceptfds` specifies a set of file descriptors on which to wait, until one or more gets an exceptional condition (arrival of out-of-band data on a network condition or certain conditions on a terminal in *packet mode*).



# **select ( ) Time Specifications**

- If `t_ptr` is NULL, then block forever, until some descriptor becomes ready.
- If elements of `*t_ptr` are 0, then return immediately after testing descriptors.
- If `t_ptr` is non-NULL and elements are non-zero, then return after specified time. If `readfds`, `writelfds` and `exceptfds`, this gives a sleep with finer granularity than `sleep ( )`.

# Operating on File Descriptor Sets

A `fd_set` is a set of file descriptors. Can only be operated on using the following macros:

`FD_ZERO(fd_set *fdset)`

Clear entire `fdset`.

`FD_SET(int fd, fd_set *fdset)`

Set `fd` in set `fdset`.

`FD_CLR(int fd, fd_set *fdset)`

Clear `fd` in set `fdset`.

`FD_ISSET(int fd, fd_set *fdset)`

Return non-zero if descriptor `fd` is set in `fdset`.

# Select Example

```
int
main(int argc, char *argv[])
{
    int fifoFD;
    if (argc != 2) {
        fprintf(stderr, "usage: nonblock FIFO\n");
        exit(1);
    }
    if ((fifoFD = open(argv[1], O_RDONLY)) < 0) {
        fprintf(stderr, "could not open %s: %s\n",
                argv[1], strerror(errno));
        exit(1);
    }
}
```

# Select Example Continued

```
while (1) {
    fd_set readfds;
    FD_ZERO(&readfds);
    FD_SET(STDIN_FILENO, &readfds);
    FD_SET(fifoFD, &readfds);
    if (select(fifoFD + 1, &readfds,
               NULL, NULL, NULL) < 0) {
        fprintf(stderr, "select error: %s\n",
                strerror(errno));
        exit(1);
    }
    else {
        int fd = (FD_ISSET(fifoFD, &readfds))
            ? fifoFD
            : STDIN_FILENO;
        char *src = (fd == fifoFD)
            ? "FIFO"
            : "STDIN";
```

# Select Example Continued

```
enum { MAX_BUF = 128 };
char buf[MAX_BUF];
int n;
if ((n = read(fd, buf, MAX_BUF))
    < 0) {
    fprintf(stderr,
            "read error on %s: %s\n",
            src, strerror(errno));
    exit(1);
}
else if (n == 0) {
    return 0;
}
else {
    printf("%s: %.*s", src, n, buf);
}
}
}
```

# Controlling an Open File

File locking can be done using `fcntl()` or `flock()` (not always available).

Recall:

```
int fcntl(int fd, int op, ...);
```

- `fd` is descriptor of already opened file.
- `op` describes operation which may use the optional 3rd argument: Covered `F_DUPFD`, `F_GETFD`, etc. Will look at `F_GETLK`, `F_SETLK`, `F_SETLKW`.

# File Locking Using `fcntl()`

When using `F_GETLK`, `F_SETLK`, `F_SETLKW`, 3rd argument specified as:

```
struct flock {
    short l_type; /* lock type: one of F_RDLCK, F_WRLCK, F_UNLCK. */
    short l_whence; /* interpretation for l_start: as for lseek() */
    off_t l_start; /* start of locked section */
    off_t l_len; /* length of locked section: 0 means to EOF from l_start */
    pid_t l_pid; /* F_GETLK returns PID of process holding lock. */
};
```

Typically many processes can concurrently *read* a resource, while only a single process can *write* the resource. Hence *read-locks* are also referred to as *shared locks* while *write-locks* are referred to as *exclusive locks*.

`F_SETLK` returns with `errno` set to `EAGAIN` or `EACCES` if lock cannot be obtained, whereas `F_SETLKW` will block.

# Lock Inheritance

Locks are associated with a process and a file.

- When process terminates all its locks are released.
- When a descriptor is closed, all locks on file associated with descriptor are released. This is true even if the file is open under another descriptor in the same process.
- Locks are not inherited by a child across a `fork()`.
- Locks are inherited across an `exec()`. However, if the close-on-exec flag is set for a file descriptor, then all locks on the underlying file are released when the descriptor is closed as part of the `exec()`.



# Advisory versus Mandatory Locking

- With *advisory locking* locking will work only if processes cooperate.
- On some systems, *mandatory locking* can be obtained if the `setgid` bit is set and group-execute bit is off.
- On Linux, filesystem has to be mounted with `-o mand` option.
- If a mandatory exclusive lock is hit, blocking read/write will block, non-blocking read/write get an error (`EAGAIN`).
- If a `open( )` tries to truncate (`O_TRUNC`) an existing file with outstanding mandatory locks, then error `EAGAIN` (solaris gives error even for `O_CREAT`).
- Mandatory locking works at system call level. Can have surprising effects at application level in that it is possible to *edit* a file with pending

mandatory locks!!

# Excluding Multiple Processes

Following program:

```
static int
lockPidFile()
{
    int fd = open(FILE_NAME, O_WRONLY|O_CREAT, 0644);
    if (fd < 0) {
        fprintf(stderr, "cannot open %s: %s\n",
                FILE_NAME, strerror(errno));
        exit(1);
    }
    else {
        struct flock flock;
        memset(&flock, 0, sizeof(flock));
        flock.l_type = F_WRLCK;
        flock.l_whence = SEEK_SET;
        flock.l_start = 0;
        flock.l_len = 0;
```

# Excluding Multiple Processes Continued

```
if (fcntl(fd, F_SETLK, &flock) < 0) {
    if (errno == EAGAIN &&
        fcntl(fd, F_GETLK, &flock) >= 0) {
        fprintf(stderr, "process %ld has lock\n",
            (long)flock.l_pid);
    }
    else {
        fprintf(stderr, "lock error: %s\n",
            strerror(errno));
    }
    exit(1);
}
if (ftruncate(fd, 0) != 0) {
    fprintf(stderr, "cannot truncate %s: %s\n",
        FILE_NAME, strerror(errno));
    exit(1);
}
```

# Excluding Multiple Processes Continued

```
{
    FILE *f = fdopen(fd, "w");
    if (!f) {
        fprintf(stderr, "cannot fdopen %s(%d): %s\n",
            FILE_NAME, fd, strerror(errno));
        exit(1);
    }
    fprintf(f, "%ld\n", (long)getpid());
    fflush(f);
}
}
return 0;
}
```

# Excluding Multiple Processes Continued

```
int main() {  
    printf("%ld trying lock\n", (long)getpid());  
    lockPidFile();  
    printf("%ld got lock\n", (long)getpid());  
    sleep(60);  
    return 0;  
}
```

# Excluding Multiple Processes Log

```
$ ./fcntl_lock &
[1] 30598
$ 30598 trying lock
30598 got lock

$ ./fcntl_lock &
30599 trying lock
[2] 30599
$ process 30598 has lock

[2]+  Exit 1                      ./fcntl_lock
$ ps
  PID TTY          TIME CMD
 30473 pts/15      00:00:00 bash
 30598 pts/15      00:00:00 fcntl_lock
 30603 pts/15      00:00:00 ps
$
[1]+  Done                      ./fcntl_lock
$ ./fcntl_lock &
[1] 30626
30626 trying lock
30626 got lock
$
```

# Memory Mapped I/O

- Allow accessing files using normal memory access operations instead of `read()`, `write()`, `seek()`, etc.
- Can also be used for IPC between arbitrary processes.
- Often used by many `mmap` implementations for allocating requests for large memory blocks. Also used by GNU's `mmap` library.



# Mapping Types

## **Private file mapping**

Maps a region of a file into memory, but writes to the memory by one process are not visible to other processes (uses copy-on-write).

## **Shared file mapping**

Maps a region of a file into memory and modifications to the memory are reflected in the file and visible to other processes (does not use copy-on-write).

## **Private anonymous mapping**

No corresponding file. Uses copy-on-write so that other processes do not see changes. Used for allocating large blocks of zero-filled memory.

## **Shared anonymous mapping**

No corresponding file. No copy-on-write. Can be used for sharing memory between multiple related processes.

# Memory Mapped Files

```
void *mmap(void *addr, size_t len, int prot,  
           int flags, int filedес, off_t off);
```

`addr`

Specifies address at which `filedes` should be mapped. Normally specified as 0, to allow system to choose address. If non-zero, should be a multiple of the system page size, especially if `MAP_FIXED` flag is specified.

`len`

Number of bytes to be mapped.

`prot`

Must be specified as consistent combinations (or) of `PROT_READ`, `PROT_WRITE`, `PROT_EXEC` or `PROT_NONE`. Specified value must be consistent with mode on `filedes`. Because of hardware limitations, actual permissions may be less restrictive than those requested.

# Memory Mapped I/O

## Continued

`flags`

Consistent combinations of

`MAP_FIXED`

Map at specified address `addr` only.

`MAP_ANON`

Create an anonymous mapping not connected to a file. `fd` and `off` are ignored. Mapped memory is initialized with zeros. On some systems, where `MAP_ANON` is not supported, the same result may be obtained by mapping `/dev/zero` (further details below).

`MAP_SHARED`

Any writes to the memory update the file.

`MAP_PRIVATE`

Any writes to the memory update a copy of the file.

One of `MAP_SHARED` or `MAP_PRIVATE`  
must be specified.

`filedes`

File descriptor being mapped.

`off`

Starting offset of region in `filedes` being  
mapped.

# Unmapping Memory

```
int munmap(void *addr, size_t len);
```

- Closing `filedes` does not unmap.
- Unmapped when process terminates, or by calling `unmap( )`.

# Memory Mapped File Copy

Timing results inconclusive.

```
int
main(int argc, char *argv[])
{
    int fdin, fdout;
    char *src, *dst;
    struct stat statbuf;

    if (argc != 3) {
        fprintf(stderr,
                "usage: a.out <fromfile> <tofile>\n");
        exit(1);
    }
    if ( (fdin = open(argv[1], O_RDONLY)) < 0) {
        fprintf(stderr,
                "can't open %s for reading: %s\n",
                argv[1], strerror(errno));
        exit(1);
    }
    if ( (fdout =
        open(argv[2], O_RDWR|O_CREAT|O_TRUNC,
            0644)) < 0) {
        fprintf(stderr,
                "can't creat %s for writing: %s\n",
                argv[1], strerror(errno));
        exit(1);
    }
}
```

# Memory Mapped File Copy

## Continued

```
if (fstat(fdin, &statbuf) < 0) {
    /* need size of input file */
    perror("fstat error"); exit(1);
}
/* set size of output file */
if (lseek(fdout, statbuf.st_size - 1, SEEK_SET)
    == -1) {
    perror("lseek error"); exit(1);
}
if (write(fdout, "", 1) != 1) {
    perror("write error"); exit(1);
}

if ( (src = mmap(0, statbuf.st_size, PROT_READ,
                 MAP_FILE | MAP_SHARED, fdin, 0))
     == (caddr_t) -1) {
    perror("mmap error for input"); exit(1);
}
if ((dst=
     mmap(0, statbuf.st_size, PROT_READ|PROT_WRITE,
          MAP_FILE | MAP_SHARED, fdout, 0))
    == (caddr_t) -1) {
    perror("mmap error for output"); exit(1);
}

/* do the file copy */
memcpy(dst, src, statbuf.st_size);

exit(0);
}
```

# Shared Memory IPC

- The special device file `‘/dev/zero’` is an infinite source of 0 bytes when read.
- Ignores any writes to it.
- If `/dev/zero` is mapped in a parent, using `MAP_SHARED`, then writes to the corresponding mapped memory are seen by all descendent processes.
- Allows IPC between related processes.
- Still need some form of synchronization.



# **mprotect ( ) Call**

```
int mprotect(const void *addr,  
             size_t len, int prot);
```

- Usually `addr` must be page-aligned `len` should be a multiple of the page-size (use `getpagesize( )`) (otherwise it is rounded up to a page size boundary).
- `prot` is a bitwise-or of `PROT_NONE`, `PROT_READ`, `PROT_WRITE`, `PROT_EXEC`.

# Auto Stack Growing

- Translate infix expressions to postfix using a recursive descent parser (details unimportant; covered previously).
- Evaluate postfix expressions using a stack with size doubled on overflow.
- Use `mprotect()` to turn off access above stack. Catch `SIGSEGV` signal and double stack size. Note that catching this signal and continuing may not work on all systems.
- Need to restart instruction which blew stack: use `sigsetjmp()`, `siglongjmp()`.

# Execution Log

```
$ ./mmapplay
>> 1+2
postfix: 1 2+
interpreted result (no stk check): 3
interpreted result (with stk check): 3
>> 1+2*3
postfix: 1 2 3*+
interpreted result (no stk check): 7
stack grown to 4 ints
interpreted result (with stk check): 7
>> 1 + (2 + (3 + (4 + (5 + (6 + (7 + (8 + 9)))))))
postfix: 1 2 3 4 5 6 7 8 9+++++++
interpreted result (no stk check): 45
stack grown to 4 ints
stack grown to 8 ints
stack grown to 16 ints
interpreted result (with stk check): 45
>> ^D$
```

# main( ) Program

```
int
main()
{
    enum {
        IN_BUF_LEN = 80,
        OUT_BUF_LEN = 80
    };
    char inBuf[IN_BUF_LEN];
    char outBuf[OUT_BUF_LEN];
    while (printf(">> ") && fflush(stdout) == 0 &&
           fgets(inBuf, IN_BUF_LEN, stdin)) {
        int parseStatus;
        if (inBuf[strlen(inBuf) - 1] != '\n') {
            fprintf(stderr, "input too long\n"); continue;
        }
        parseStatus = parse(inBuf, outBuf, OUT_BUF_LEN);
    }
```

# main( ) Program Continued

```
if (parseStatus == -1) {
    fprintf(stderr, "parse error\n");
}
else if (parseStatus == -2) {
    fprintf(stderr, "output buffer overflow\n");
}
else {
    fprintf(stdout, "postfix: %s\n", outBuf);
    { int sp[16];
      int result = interp(outBuf, sp);
      fprintf(stdout,
              "interpreted result (no stk check): %d\n",
              result);
    }
    {
        int result;
        initStk(2);
        result = interp(outBuf, NULL);
        freeStk();
        fprintf(stdout,
                "interpreted result (with stk check): %d\n",
                result);
    }
}
return 0;
}
```

# Stack Initialization

```
static void *mapBase;
int mapLen;

int *sp;
static int *stkBase;

void
initStk(int size)
{
    size_t ps = getpagesize();
    int zero = open("/dev/zero", O_RDWR);
    mapLen = 2*ps;
    if (zero < 0) {
        perror("/dev/zero open"); exit(1);
    }
    mapBase = mmap(NULL, mapLen,
                   PROT_READ|PROT_WRITE,
                   MAP_PRIVATE, zero, 0);
    if (mapBase == (void *)-1) {
        perror("mmap"); exit(1);
    }
    assert(((long)mapBase) % ps == 0);
}
```

# Stack Initialization Continued

```
{ char *protPage = (char *)mapBase + ps;
  if (mprotect(protPage, ps, PROT_NONE) < 0) {
    perror("mprotect"); exit(1);
  }
}
{ int *protPage = (int *)((char *)mapBase + ps);
  sp = stkBase = protPage - size;
}
{ struct sigaction sigact;
  sigset_t sigmask;
  sigemptyset(&sigmask);
  sigact.sa_handler = growStk;
  sigact.sa_mask = sigmask;
  sigact.sa_flags = 0;
  if (sigaction(SIGSEGV, &sigact, NULL) < 0) {
    perror("signal"); exit(1);
  }
}
}
```

# Stack Deallocation

```
void
freeStk(void)
{
    if (munmap(mapBase, mapLen) < 0) {
        perror("munmap"); exit(1);
    }
}
```



# Stack Overflow Handler

```
sigjmp_buf stkJump;

/* double size of stack */
void
growStk(int signo)
{
    int initStkSize = sp - stkBase;
    int *newBase = stkBase - initStkSize;
    int i;
    for (i = 0; i < initStkSize; i++) {
        newBase[i] = stkBase[i];
    }
    sp = stkBase; stkBase = newBase;
    fprintf(stderr, "stack grown to %d ints\n",
            2*initStkSize);
    siglongjmp(stkJump, 1);
}
```

# Postfix Interpreter

```
int
interp(const char *postfix, int *sp1)
{
    const char *p = postfix;
    if (sp1 != NULL) sp = sp1;
    while (*p != '\0') {
        switch (*p++) {
            case ' ': {
                char *pp;
                volatile int v = strtol(p, &pp, 10);
                p = pp;
                sigsetjmp(stkJump, 1);
                *sp = v;
                sp++;
            }
            break;
            case '+': {
                int right = *--sp;
                int left = *--sp;
                *sp++ = left + right;
            }
            break;
        }
    }
}
```

# Postfix Interpreter

## Continued

```
case '-': {
    int right = *--sp;
    int left = *--sp;
    *sp++ = left - right;
}
break;
/* Code for * / similar */
case 'u': {
    int rand = *--sp;
    *sp++ = -rand;
}
break;
}
}
return *--sp;
}
```

# References

Text, 5.9, 13, 44.9, 49, 55, 63

APUE: Ch. 14.