

Signals

- Common Signals
- Common Signals Continued
- Signals Continued
- Signal Function
- Sending Signals
- Sending Signals Continued
- An Example Signal Handler
- Example Signal Handler Execution Log
- Signal Gotchas
- Unreliable Signals
- Remembering Occurrence of a Signal
- Interrupted System Calls
- Interrupted System Calls Continued
- Non-Reentrant Functions
- SIGCHLD versus SIGCLD
- Implement `sleep()` Using `alarm()`: First attempt
- Problems with `sleep()` First Attempt
- Implement `sleep()` Using `alarm()`: Second attempt
- `sleep()` Second Attempt: Review
- Reliable Signal Terminology

- Signal Sets and `sigprocmask()`
- `sigpending()`
- `sigaction()`
- `sigaction()` Flags
- `sigaction()` Flags Continued
- Reliable `signal()`
- `sigsetjmp()` and `siglongjmp()`
- Preventing a Signal in a Critical Region
- `sigsuspend()` Function
- A Final `sleep()`
- A Final `sleep()` Continued
- Real-Time Signals
- Real-Time Signal Details
- Sending Real-Time Signals
- References

Common Signals

SIGABRT

Generated by calling `abort()`.

SIGALRM, SIGPROF, SIGVTALRM

Generated by calling `alarm()` or `setitimer()`.

SIGBUS, SIGEMT, SIGFPE, SIGIOT, SIGILL,
SIGPWR, SIGSEGV, SIGTRAP

Indicates hardware conditions.

SIGCHLD

Indicates a child termination.

SIGCONT, SIGSTOP, SIGTSTP, SIGTTIN, SIGTTOU

Job control signals.

SIGHUP

Indicates terminal hangup. Sent to session leader. If session leader terminates, then it is sent to each foreground process.

Common Signals Continued

SIGINFO

Information requested on foreground processes (^T).

SIGINT

Signal generated to foreground processes by interrupt char (^C).

SIGIO, SIGPOLL

I/O signals.

SIGKILL

Kills a process. Cannot be caught.

SIGPIPE

Caused by a write to a pipe not being read.
Similarly for sockets.

SIGQUIT

Terminates foreground process group with core dump.

Signals Continued

SIGSYS

Bad system call.

SIGTERM

Terminate process.

SIGURG

Urgent condition. Used for out-of-band data on network.

SIGUSR1, SIGUSR2

Available for any purpose by user.

SIGWINCH

Generated if window size changes.

SIGXCPU, SIGXFSZ

Generated if soft limit exceeded.

Signal Function

```
typedef void SigFunc(int);  
SigFunc *signal(int signo, SigFunc *handler);
```

or without the typedef:

```
void (*signal(int signo, void (*handler)(int)))(int);
```

- `signo` specifies signal whose disposition is being changed.
- `handler` specifies the *signal handler*.
- Returns previous signal disposition if ok, `SIG_ERR` on error.
- No way to query signal disposition without changing it.
- Special dispositions of `SIG_IGN` (ignore the signal) or `SIG_DFL` (set the default disposition).
- In original Unix API with various problems and incompatibilities; modern programs should use `sigaction()`

Sending Signals

```
int kill(pid_t pid, int signo);
```

```
int raise(int signo);
```

- `raise()` allows a process to send a signal to itself (ANSI-C).
- Operation of `kill` depends on `pid`:

```
pid > 0
```

Signal sent to specified process.

```
pid == 0
```

Signal sent to all process with same group ID as sending process.

```
pid < 0
```

Sent to all processes with group ID equal to absolute value of `pid`.

```
pid == -1
```

Used for broadcasting.

Sending Signals Continued

- Super user can send signal to any process.
- Non super users can only send signals to processes with same real or effective UID (with exception of `SIGCONT` which can be sent by any process in same session).
- `kill(1)` is shell interface to `kill(2)`.
- `signo == 0` can be used to query existence of a process.

An Example Signal Handler

```
int
main(void)
{
    if (signal(SIGUSR1, sig_usr) == SIG_ERR) {
        perror("SIGUSR1"); exit(1);
    }
    if (signal(SIGUSR2, sig_usr) == SIG_ERR) {
        perror("SIGUSR2"); exit(1);
    }
    for ( ; ; )
        pause();
}

static void
sig_usr(int signo) /* argument is signal number */
{
    if (signo == SIGUSR1) {
        printf("received SIGUSR1\n"); /* unsafe */
    }
    else if (signo == SIGUSR2) {
        printf("received SIGUSR2\n"); /* unsafe */
    }
    else {
        fprintf(stderr, "received signal %d\n", signo);
        exit(1);
    }
    return;
}
```

Example Signal Handler Execution Log

```
$ ./sigusr &  
[1] 4348  
$ kill -USR1 4348  
$ received SIGUSR1
```

```
$ kill -USR2 4348  
$ received SIGUSR2
```

```
$ kill 4348  
$  
[1]+  Terminated                  ./sigusr  
$
```

Signal Gotchas

- Unreliable signals.
- Non-reentrant library functions.
- Inconsistent semantics (`SIGCLD` and `SIGCHLD`).
- Slow system calls may or may not be automatically restartable when interrupted by a signal.
- Race conditions.

Unreliable Signals

- Signal disposition reset to default when the signal occurred.
- Usually remedied by reinitializing the handler within the handler:

```
void sigInt(int signo) {  
    if (signal(SIGINT, sigInt)) { ... }  
    ...  
}
```

- However, if signal occurred again before handler was reinitialized, then default action for signal would occur (usually process termination).
- Program would appear to work most of the time.

Remembering Occurrence of a Signal

- Use a flag to remember occurrence of a signal:

```
int sigIntFlag = 0;
...
signal(SIGINT, sigInt);
...

while (sigIntFlag == 0) pause;
...
void sigInt(int signo) {
    signal(SIGINT, sigInt);
    sigIntFlag = 1;
}
```

- What if signal occurs between test of sigIntFlag and pause()?

Interrupted System Calls

- *Slow* devices are those which can block *forever*: includes pipes, terminal devices, network devices, some IPC and `ioctl()` operations, as well as `pause()`.
- If a signal occurs while a system call is accessing a slow device, then in early Unix systems, the call returned with `errno` set to `EINTR`.

```
again:
    if ((n = read(fd, buf, BUFSIZE)) < 0) {
        if (errno == EINTR) goto again;
        ...
    }
    ...
```

Interrupted System Calls Continued

- 4.2 BSD automatically restarted such system calls which were interrupted. Consider

```
alarm(TIME_OUT);  
if ((n = read(fd, buf BUFSIZE)) < 0) {  
    if (timeOutExpired) { ... }  
    ...  
}  
...
```

Above code which attempts to put a time-out on a read, will not work with above semantics (it also has a race condition).

- 4.3 BSD allowed control over which signals interrupt systems calls and which signals lead to the call being automatically restarted.
- SysV never automatically restarted interrupted system calls by default.

Non-Reentrant Functions

- Non-reentrant functions cannot be called from a signal handler.
- POSIX specifies that certain functions are *async-signal-safe* (listed in text). Functions not on the list should not be called from signal handler.
- It is usually not safe to call any of the C library functions (like `printf()` or `malloc()`) unless the signal handler will also terminate the program, or the library has been guaranteed to be reentrant.
- Only 1 `errno` per process; hence it is necessary to save and restore `errno` within signal handler.

SIGCHLD versus SIGCLD

- Semantics of BSD `SIGCHLD` are normal, like that of any other signal.
- System V has unusual semantics for `SIGCLD`:
 - If disposition explicitly set to `SIG_IGN`, then no zombies.
 - If disposition set to be handled, then kernel checks to see if there are any terminated children, in which case it immediately calls the handler.
- Consider the need to reestablish the signal handler when the handler is first entered ... leads to a recursive loop!!

Implement sleep() Using alarm() : First attempt

```
static void
sigAlarm() {
    return;
}

unsigned int
sleep1(unsigned int nSecs) {
    if (signal(SIGALRM, sigAlarm)
        == SIG_ERR) {
        return nSecs;
    }
    alarm(nSecs);
    pause();
    return alarm(0);
}
```

Problems with `sleep()`

First Attempt

1. If caller of `sleep1()` has a `alarm()` set, then that `alarm()` is lost. Can be corrected by using return value of `alarm()`.
2. Disposition of `SIGALRM` changed. Can be fixed by saving return value of `signal()` and restoring it before returning.
3. Race condition between `alarm()` and `pause()`.

Implement sleep() Using alarm() : Second attempt

```
static jmp_buf alarmEnv;

static void
sigAlarm(int signo)
{
    longjmp(alarmEnv, 1);
}

unsigned int
sleep2(unsigned int nSecs)
{
    if (signal(SIGALRM, sigAlarm) == SIG_ERR) {
        return nSecs;
    }
    if (setjmp(alarmEnv) == 0) {
        alarm(nSecs);
        pause();
    }
    return alarm(0);
}
```

`sleep()` Second Attempt: Review

- Fixes problem (3).
- Assume that (1) and (2) can be taken care of.
- If alarm interrupts another signal handler, then that signal handler is aborted!!

Reliable Signal Terminology

- A signal is *delivered* to a process when the action for a signal is taken.
- A signal is *pending* during the time it is generated and the time it is delivered.
- With reliable signals, a process has the option of *blocking* a signal. If a blocked signal occurs and its disposition is default or catch, then the signal remains pending until the process unblocks it or sets its disposition to ignore.
- The delivery of a blocked signal depends on its disposition at the time it is delivered, not the time at which it was generated.
- If more than 1 occurrence of the same signal can be pending, then the system may *queue* the signals. Most systems do not queue signals and only a single signal will be delivered.

Signal Sets and **sigprocmask()**

```
int sigemptyset(sigset_t *set);
int sigfillset(sigset_t *set);
int sigaddset(sigset_t *set, int sigNum);
int sigdelset(sigset_t *set, int sigNum);
int sigismember(const sigset_t *set, int sigNum);

int sigprocmask(int how, const sigset_t *set1,
                sigset_t *oset);
```

- how is used when set1 is non-null: SIG_BLOCK blocks set1 signals; SIG_UNBLOCK unblocks set1 signals; SIG_SETMASK sets signal mask to set1.
- If oset non-NULL, then previous signal mask is returned via oset.

sigpending()

```
int sigpending(sigset_t *set);
```

- `sigpending()` returns all pending signals. Use `sigismember()` to process return value.

sigaction()

```
struct sigaction {  
    void (*sa_handler)();  
    sigset_t sa_mask;  
    int sa_flags;  
} SigAction;  
int sigaction(int signo,  
              const SigAction *act,  
              SigAction *oact);
```

- `sa_handler` field of `SigAction` allows specifying a handler just as for `signal()`.
- `sa_mask` indicates signals which should be masked out during handler execution. When handler returns, old mask is automatically restored.
- When handler is invoked, `signo` is always added to the mask which is installed before the handler is entered.

sigaction() Flags

`sa_flags` member of `SigAction` allows `sigaction()` to emulate different behaviors of `signal()`:

`SA_NOCHLDSTOP`

When `signo` is `SIGCHLD`, do not generate signal when a child stops (still generated when a child terminates).

`SA_RESTART`

Systems calls interrupted by this signal are automatically restarted.

`SA_ONSTACK`

If an alternate stack defined using `sigaltstack()` use alternate stack for delivery of this signal.

`SA_NOCLDWAIT`

Emulate Sys V behavior of not creating zombies for terminated children.

sigaction() Flags

Continued

SA_NODEFER

Signal not automatically blocked when handler is entered.

SA_RESETHAND

Reset the handler to `SIG_DFL` before handler is entered.

SA_SIGINFO

Pass additional information (2 additional arguments) to signal handler.

Simulate unreliable signals using

`SA_NODEFER | SA_RESETHAND`.

Reliable signal()

```
typedef void (SigFunc)(int sigNum);

SigFunc *
signal(int sigNum, SigFunc *func)
{
    struct sigaction act, oact;

    act.sa_handler = func;
    sigemptyset(&act.sa_mask);
    act.sa_flags = 0;
    if (signo == SIGALRM) {
#ifdef SA_INTERRUPT
        act.sa_flags |= SA_INTERRUPT; /*Sun OS */
#endif
    }
    else {
#ifdef SA_RESTART
        act.sa_flags |= SA_RESTART;
#endif
    }
    if (sigaction(sigNum, &act, &oact) < 0) {
        return SIG_ERR;
    }
    return oact.sa_handler;
}
```

`sigsetjmp()` and `siglongjmp()`

```
int sigsetjmp(sigjmp_buf env, int savemask);  
int siglongjmp(sigjmp_buf env, int value);
```

- With reliable signal semantics, old signal mask is restored when a signal handler returns.
- If signal handler exits using `longjmp()`, then POSIX does not define whether or not old signal mask should be restored.
- If `sigsetjmp()` called with `savemask` non-zero, then a **corresponding** `siglongjmp()` out of a signal handler will restore the old signal mask.

Preventing a Signal in a Critical Region

```
sigset_t newMask, oldMask;

sigemptyset(&newMask);
sigaddset(&newMask, SIGINT);
if (sigprocmask(SIG_BLOCK, &newMask, &oldMask) < 0) {
    ...
}
/* critical region */
if (sigprocmask(SIG_SETMASK, &oldMask, NULL) < 0) {
    ...
}
pause(); /* wait for signal */
```

Race condition between second `sigprocmask()` and `pause()`.

sigsuspend() Function

```
int sigsuspend(const sigset_t *mask);
```

- Resets signal mask to `mask` and goes to sleep **atomically**.
- If a signal is caught and the signal handler returns then returns (with `errno` set to `EINTR`) and the signal mask set to its previous value.

A Final sleep()

```
static void
sigAlarm(int sigNum)
{
    return;
}

unsigned int
sleep(unsigned int nSecs)
{
    struct sigaction newAct, oldAct;
    sigset_t newMask, oldMask, suspMask;
    unsigned int unslept;

    newAct.sa_handler = sigAlarm;
    sigemptyset(&newAct.sa_mask);
    newAct.sa_flags = 0;
    sigaction(SIGALRM, &newAct, &oldAct);
```


A Final `sleep()` Continued

```
sigemptyset(&newMask);
sigaddset(&newMask, SIGALRM);
sigprocmask(SIG_BLOCK, &newMask, &oldMask);

alarm(nSecs);

suspMask = oldMask;
sigdelset(&suspMask, SIGALRM);
sigsuspend(&suspMask);

unslept = alarm(0);
sigaction(SIGALRM, oldAct, NULL);
sigprocmask(SIG_SETMASK, &oldMask, NULL);

return unslept;
}
```

Real-Time Signals

- Added by Posix.1b.
- Supported if `_POSIX_REALTIME_SIGNALS` is defined.
- Allows passing information (`int` value or `void *` pointer) to signal handler.
- Allows prioritizing of signals (lower signal numbers have higher priority).
- Allows queuing of signals.

Real-Time Signal Details

- Real-time signal numbers between SIGRTMIN to SIGRTMAX with RTSIG_MAX real-time signals in between (min. 8).
- Added a additional handler to struct sigaction:

```
struct sigaction {  
    void (*sa_handler)();           /* SIG_DFL, SIG_IGN, or  
                                    * pointer to function */  
    void (*sa_sigaction)            /* Real-time signal */  
        (int, siginfo_t *, void *); /* handler function */  
    sigset_t sa_mask;               /* additional signals to be blocked  
                                    * during execution of handler */  
    int sa_flags;                   /* special flags and options */  
};
```

- New handler has prototype:

```
void handler(int signo, siginfo_t *info, void *context);
```

context is currently undefined.

- siginfo_t contains at least the following members:

```
int si_signo;          /* signal #.  Same as signo */
int si_code;           /* one of SI_USER, SI_QUEUE, SI_TIMER,
                        * SI_ASYNCIO, SI_MESGQ */
union sigval si_value; /* union { int sival_int;
                        *         void *sival_ptr;
                        *         };
```

Sending Real-Time Signals

```
int sigqueue(pid_t pid, int signo, const union sigval value);
```

- Additional parameter specifies information sent to the handler via the `info` argument.
- To guarantee queuing, `SA_SIGINFO` must be set in `sa_flags` in `struct sigaction`.
- Multiple signals with the same number generated by `sigqueue()` are queued upto a max of `SIGQUEUE_MAX` (typically 32).

References

Text: Chs 20 - 22.

APUE, Ch. 10.

Jim Frost, *UNIX Signals and Process Groups*, at
<http://www.cs.ucsb.edu/~almeroth/classes/W99.276/assignment1/signals.html>.

FSF, *The GNU C Library*, at
<http://www.gnu.org/software/libc/manual/>.